

RISC-V

Robert Baruch
robert.c.baruch@gmail.com

November 8, 2018

Chapter 1

Harts and the FENCE instruction

A hart is a hardware thread. This is in contrast to a software thread. A hart is an independent core with its own set of registers and its own program counter. However, all cores share the entire memory space.

Each hart executes its own instructions one by one, in order. However, because of caches and other accelerator magic, some other hart may not see the loads and stores from this hart in the same order! For example, imagine that hart 1 executes instructions in this order:

Store 0 to location A
Store 1 to location B
Read location A
Read location B
Add them.
Store the result to location C.

In any sane universe, 2 should be written to location C, and as far as hart 1 is concerned, that is exactly what happens.

Next, imagine that some hart is monitoring location C by reading it over and over. When it sees that location C changes, it immediate reads A and B, adds them, and compares the result with C. Maybe hart 2 is testing that hart 1 is actually adding A and B properly.

However, now imagine that hart 1 is actually accessing a fast memory cache. All the instructions are reading and writing that cache. Eventually, the cache gets flushed to main memory. However, the order that the cache writes to memory is not specified! It could write location C first, then B, then A. The cache doesn't care about the order of operations, it just cares about the view of memory according to hart 1 and the changes it made. So hart 2 might actually see 2 in C but 5 in B and 42 in A because the cache hasn't gotten around to writing locations A and B yet. The addition doesn't compare, and hart 2 thinks hart 1 is malfunctioning.

This is the essence of the *relaxed memory model*. If two harts access different areas of memory, no problem exists. But when they access the same areas of memory, things can go a bit wonky.

The FENCE instruction forces an ordering on the cache. For example, a FENCE instruction might dictate that previous writes must occur before subsequent writes. This means that a cache would be forced to write everything before the FENCE before writ-

ing anything after the FENCE. Thus, the cache could only end up writing in either of these two orders: A, B, C, or B, A, C. That is, the FENCE instruction declares that every other hart is guaranteed to see C written only after A and B are written.

So with such a FENCE, hart 2 will not see inconsistent values.

The FENCE.I instruction has to do with instruction caches and pipelines. If a hart were to write to memory which is actually code, it is possible that a cache or pipeline could still contain the old value at that address, meaning that the instruction to execute at that address is incorrect. Executing a FENCE.I instruction guarantees that the hart's caches and pipelines are updated with respect to instructions to execute.

Chapter 2

The Instructions

2.1 Abbreviations

XLEN	The processor's integer register width
imm	The unsigned immediate value in the instruction
int12	12-bit signed integer
uint12	12-bit unsigned integer
int13e	13-bit even signed integer
uint20	20-bit unsigned integer
int21e	21-bit even signed integer
uintN	$\log_2(\text{XLEN})$ -bit unsigned integer
pc	The program counter
rd	The destination register
rs1	Source register 1
rs2	Source register 2
sext(x)	x , sign-extended to XLEN bits
zext(x)	x , zero-extended to XLEN bits
(r)	The contents of a register, as a signed XLEN-bit integer
(r) _u	The contents of a register, as an unsigned XLEN-bit integer
X_b	The lowest 8 bits of X
X_h	The lowest 16 bits of X
X_w	The lowest 32 bits of X
X_d	The lowest 64 bits of X
X_n	The lowest $\log_2(\text{XLEN})$ bits of X
[a] _b	The 8-bit integer stored in memory at address a
[a] _h	The 16-bit integer stored in memory at address a
[a] _w	The 32-bit integer stored in memory at address a
[a] _d	The 64-bit integer stored in memory at address a
\leftarrow	Receives the value of
$\bar{1}$	All ones, except the least significant bit is zero
+	Addition
-	Subtraction
<<	Binary shift left

>>	Binary signed shift right
>> _u	Binary unsigned shift right
\wedge	Logical bitwise and
\vee	Logical bitwise or
\oplus	Logical bitwise exclusive or (xor)

ADD

Add (register/register)

Instruction Format: R

Opcode Type: OP

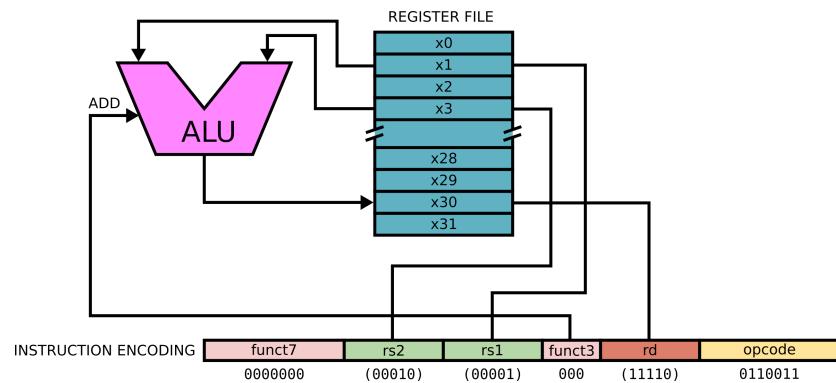
Assembly: ADD rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) + (rs2)$

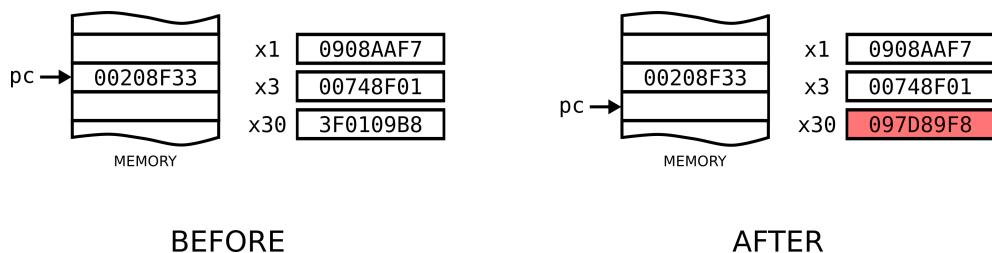
Description:

Adds the contents of register rs1 to the contents of register rs2. The result is placed in register rd.

Data Path:



Example: ADD x30, x1, x3



Notes:

Overflows of the ADD instruction for signed arithmetic can be checked using SLTI and SLT:

```
ADD  x1, x2, x3
SLTI x4, x3, 0
SLT  x5, x1, x2
BNE  x4, x5, overflow
```

ADDI

Add (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

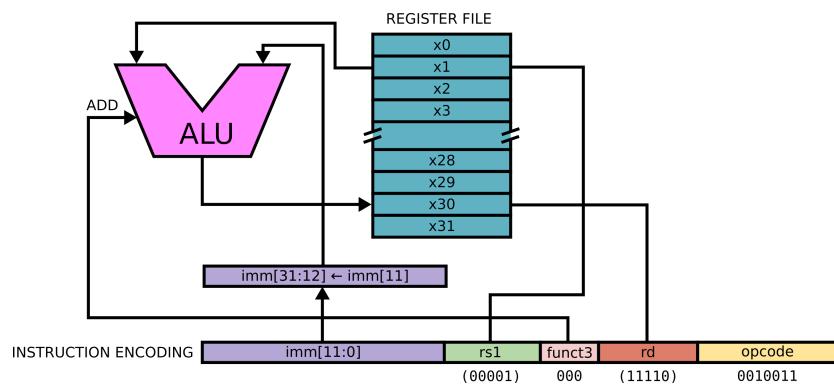
Assembly: ADDI rd, rs1, int12

Function: $(rd) \leftarrow (rs1) + \text{sext}(\text{imm})$

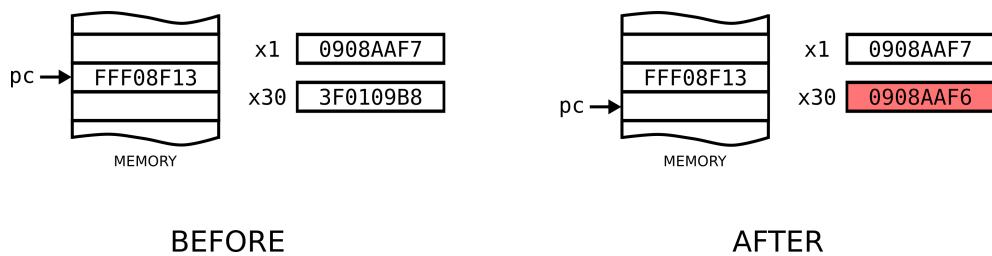
Description:

Adds the sign-extended immediate value to the contents of register rs1. The result is placed in register rd.

Data Path:



Example: ADDI x30, x1, -1



Notes:

There is no separate SUBI instruction. However, SUBI is a pseudo-instruction encoded as ADDI where the immediate value is negated. Thus, the following two instructions are encoded the same way:

```
SUBI x30, x1, 1  
ADDI x30, x1, -1
```

AND

Bitwise and (register/register)

Instruction Format: R

Opcode Type: OP

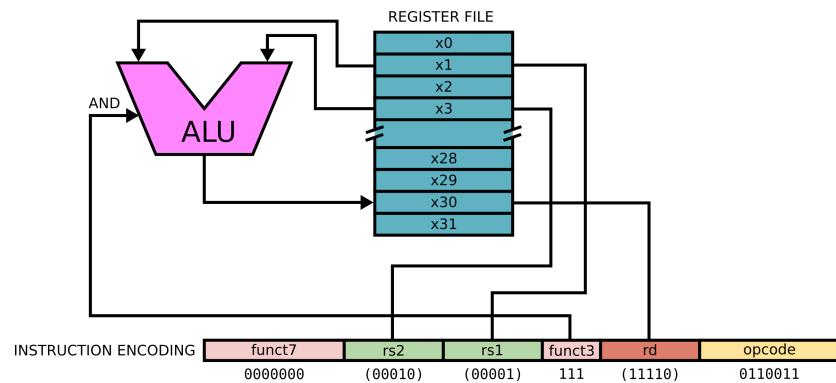
Assembly: AND rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) \wedge (rs2)$

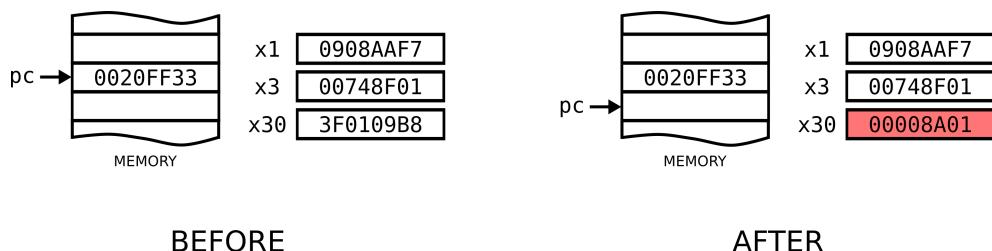
Description:

Performs a bitwise AND between the contents of register rs1 and the contents of register rs2. The result is placed in register rd.

Data Path:



Example: AND x30, x1, x3



ANDI

Bitwise and (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

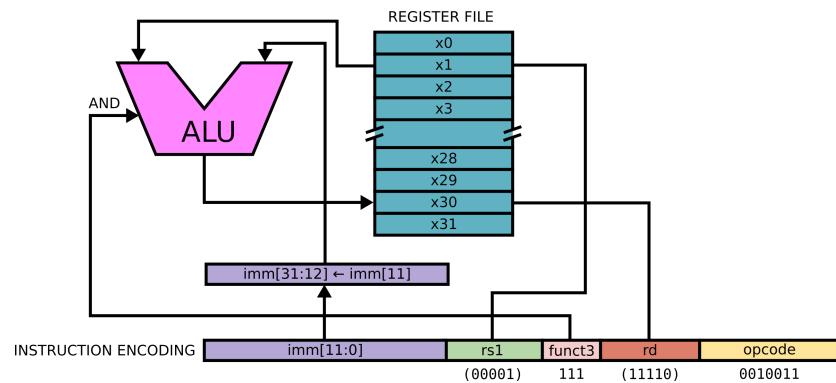
Assembly: ANDI rd, rs1, int12

Function: $(rd) \leftarrow (rs1) \wedge \text{sext(immm)}$

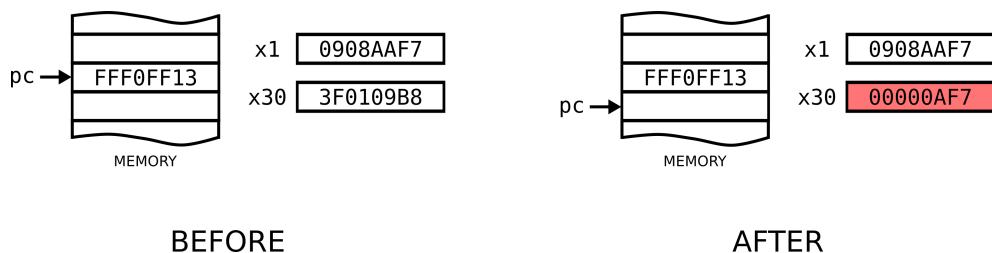
Description:

Performs a bitwise AND between the sign-extended immediate value and the contents of register rs1. The result is placed in register rd.

Data Path:



Example: ANDI x30, x1, -1



AUIPC

Add upper immediate to PC

Instruction Format: U

Opcode Type: AUIPC

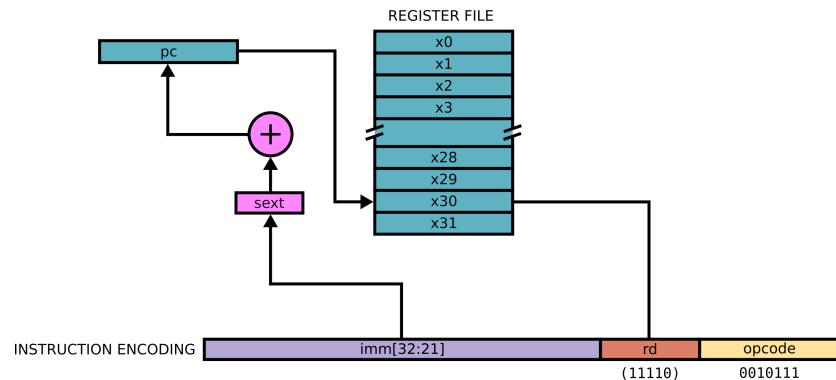
Assembly: AUIPC rd, uint20

Function: $(rd) \leftarrow (pc) + \text{sext}(\text{imm} \ll 12)$

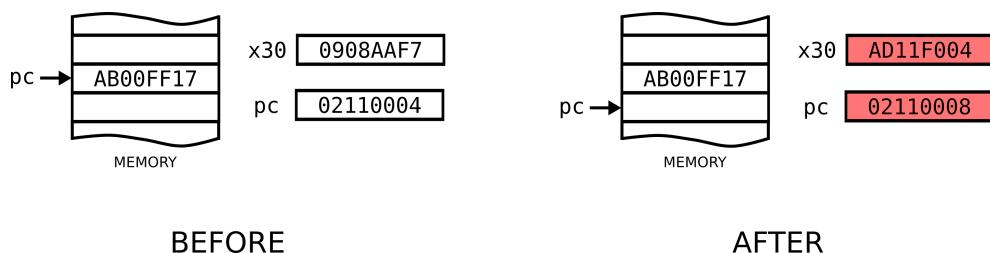
Description:

Constructs a 32-bit integer whose upper 20 bits are the immediate value, filling the lower 12 bits with zeros. The result is then sign-extended to XLEN bits, added to the program counter, and the result placed in the destination register.

Data Path:



Example: AUIPC x30, 0xAB00F



Notes:

The AUIPC instruction is used to construct pc-relative offsets outside the 12-bit range.

BEQ

Branch if equal

Instruction Format: B

Opcode Type: BRANCH

Assembly: BEQ rs1, rs2, int13e

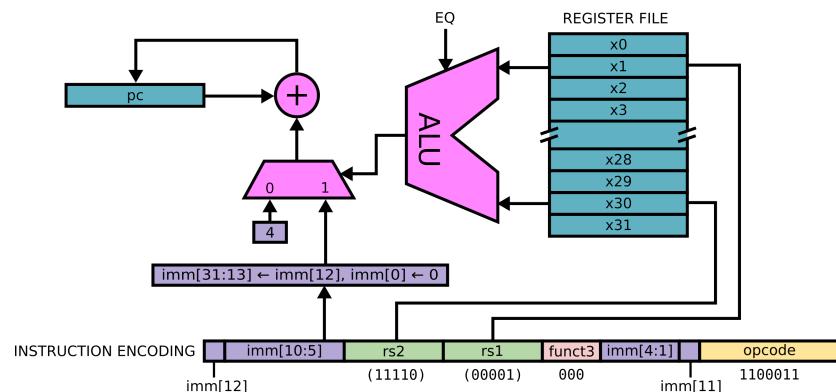
Function:

$$(pc) \leftarrow \begin{cases} (pc) + sext(imm) & \text{if } (rs1) = (rs2) \\ (pc) + 4 & \text{otherwise} \end{cases}$$

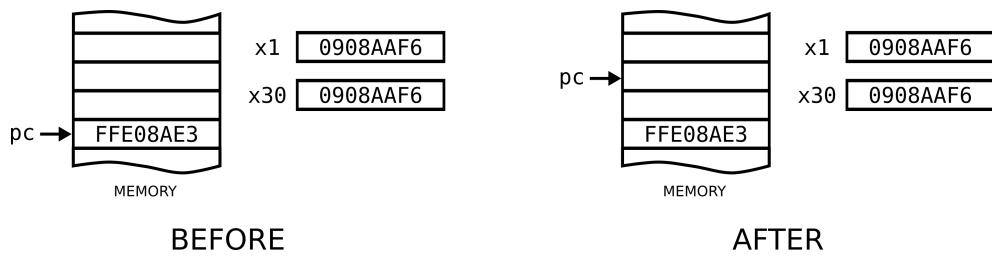
Description:

Compares the contents of register rs1 to the contents of register rs2. If they are equal, the signed offset is added to the current program counter, otherwise the program counter is set to the address of the next instruction.

Data Path:



Example: BEQ x1, x30, -8



Notes:

Offsets are limited to even addresses in the range [-4096, 4094]. Branches to addresses outside this range can be accomplished using the complementary branch followed by an unconditional jump:

```
BNE x1, x30, 8  
J    far_addr
```

A misaligned instruction fetch exception will be generated if the address of the branch taken is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

BGE

Branch if greater than or equal (signed)

Instruction Format: B

Opcode Type: BRANCH

Assembly: BGE rs1, rs2, int13e

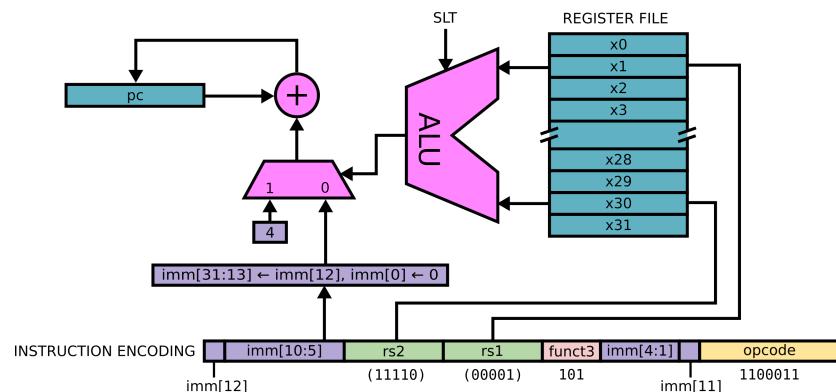
Function:

$$(pc) \leftarrow \begin{cases} (pc) + \text{sext}(imm) & \text{if } (rs1) \geq (rs2) \\ (pc) + 4 & \text{otherwise} \end{cases}$$

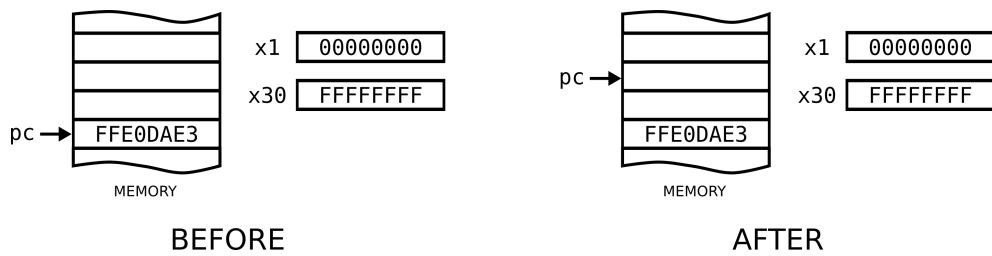
Description:

If the signed contents of register rs1 is greater than or equal to the signed contents of register rs2, the signed offset is added to the current program counter, otherwise the program counter is set to the address of the next instruction.

Data Path:



Example: BGE x1, x30, -8



Notes:

Offsets are limited to even addresses in the range [-4096, 4094]. Branches to addresses outside this range can be accomplished using the complementary branch followed by an unconditional jump:

```
BLT x1, x30, 8  
J    far_addr
```

A misaligned instruction fetch exception will be generated if the address of the branch taken is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

BGEU

Branch if greater than or equal (unsigned)

Instruction Format: B

Opcode Type: BRANCH

Assembly: BGEU rs1, rs2, int13e

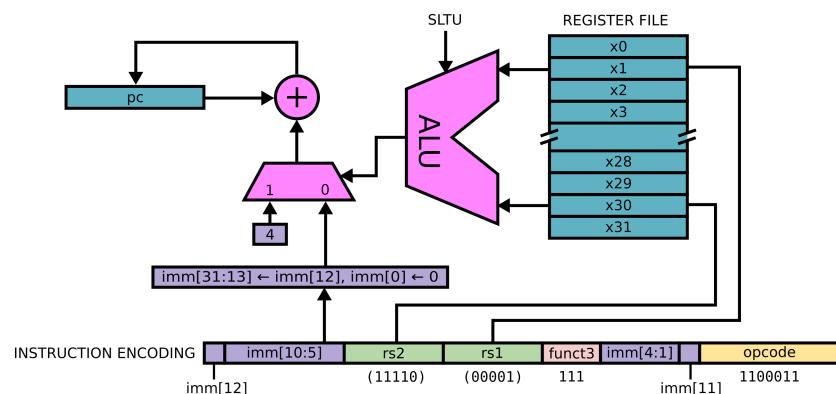
Function:

$$(pc) \leftarrow \begin{cases} (pc) + sext(imm) & \text{if } (rs1)_u \geq (rs2)_u \\ (pc) + 4 & \text{otherwise} \end{cases}$$

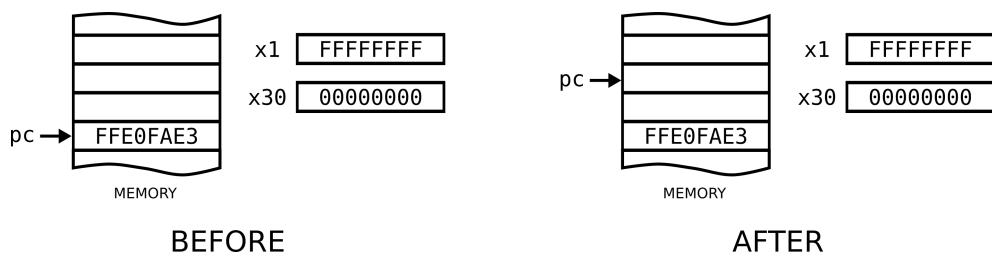
Description:

If the unsigned contents of register rs1 is greater than or equal to the unsigned contents of register rs2, the signed offset is added to the current program counter, otherwise the program counter is set to the address of the next instruction.

Data Path:



Example: BGEU x1, x30, -8



Notes:

Offsets are limited to even addresses in the range [-4096, 4094]. Branches to addresses outside this range can be accomplished using the complementary branch followed by an unconditional jump:

```
BLTU x1, x30, 8  
J      far_addr
```

A misaligned instruction fetch exception will be generated if the address of the branch taken is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

BLT

Branch if less than (signed)

Instruction Format: B

Opcode Type: BRANCH

Assembly: BLT rs1, rs2, int13e

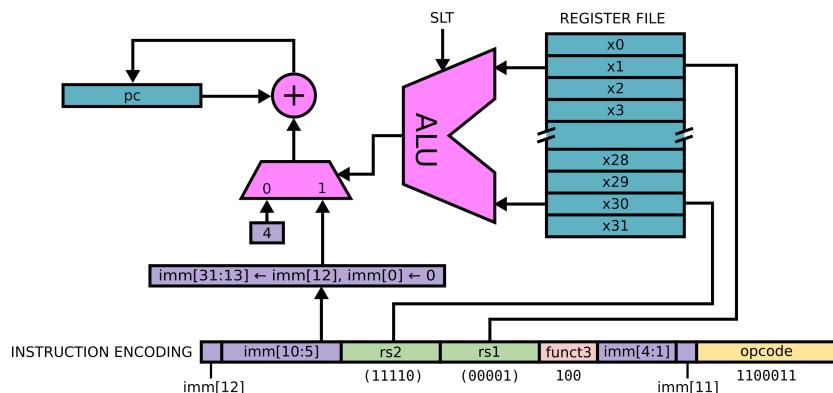
Function:

$$(pc) \leftarrow \begin{cases} (pc) + sext(imm) & \text{if } (rs1) < (rs2) \\ (pc) + 4 & \text{otherwise} \end{cases}$$

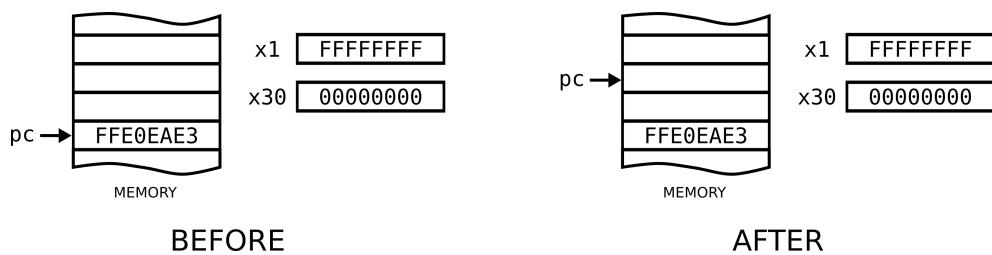
Description:

If the signed contents of register rs1 is less than the signed contents of register rs2, the signed offset is added to the current program counter, otherwise the program counter is set to the address of the next instruction.

Data Path:



Example: BLT x1, x30, -8



Notes:

Offsets are limited to even addresses in the range [-4096, 4094]. Branches to addresses outside this range can be accomplished using the complementary branch followed by an unconditional jump:

```
BGE x1, x30, 8  
J    far_addr
```

A misaligned instruction fetch exception will be generated if the address of the branch taken is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

BLTU

Branch if less than (unsigned)

Instruction Format: B

Opcode Type: BRANCH

Assembly: BLTU rs1, rs2, int13e

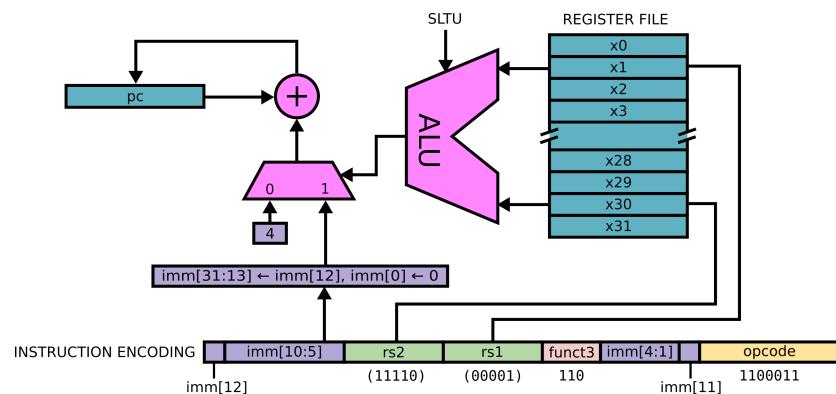
Function:

$$(pc) \leftarrow \begin{cases} (pc) + sext(imm) & \text{if } (rs1)_u < (rs2)_u \\ (pc) + 4 & \text{otherwise} \end{cases}$$

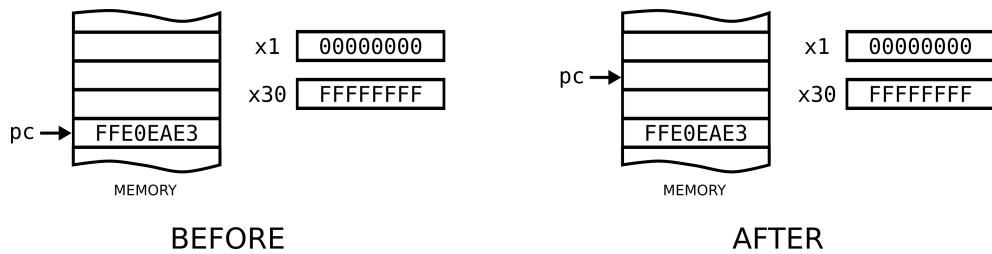
Description:

If the unsigned contents of register rs1 is less than the unsigned contents of register rs2, the signed offset is added to the current program counter, otherwise the program counter is set to the address of the next instruction.

Data Path:



Example: BLTU x1, x30, -8



Notes:

Offsets are limited to even addresses in the range [-4096, 4094]. Branches to addresses outside this range can be accomplished using the complementary branch followed by an unconditional jump:

```
BGEU x1, x30, 8  
J      far_addr
```

A misaligned instruction fetch exception will be generated if the address of the branch taken is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

BNE

Branch if not equal

Instruction Format: B

Opcode Type: BRANCH

Assembly: BNE rs1, rs2, int13e

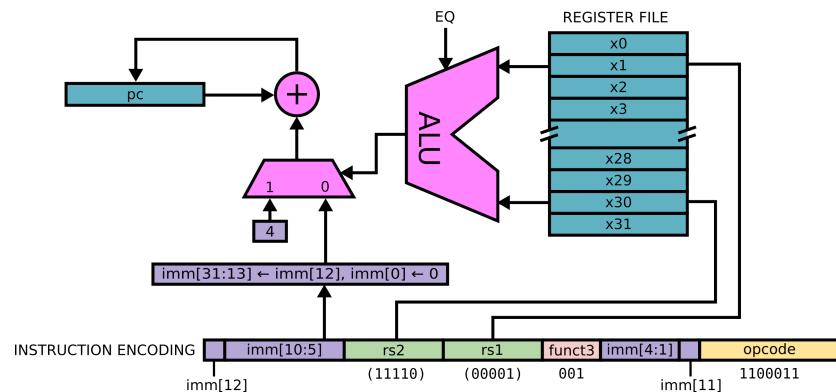
Function:

$$(pc) \leftarrow \begin{cases} (pc) + \text{sext}(\text{imm}) & \text{if } (rs1) \neq (rs2) \\ (pc) + 4 & \text{otherwise} \end{cases}$$

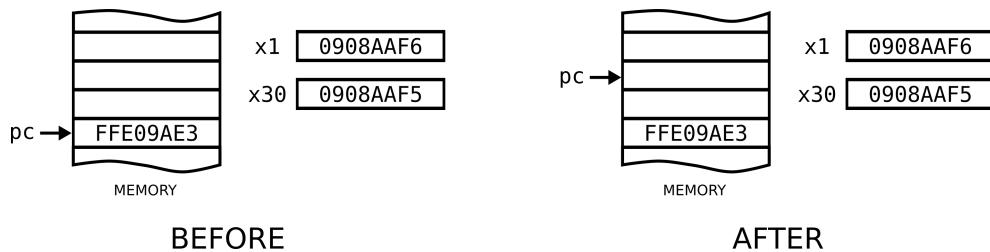
Description:

Compares the contents of register rs1 to the contents of register rs2. If they are not equal, the signed offset is added to the current program counter, otherwise the program counter is set to the address of the next instruction.

Data Path:



Example: BNE x1, x30, -8



Notes:

Offsets are limited to even addresses in the range [-4096, 4094]. Branches to addresses outside this range can be accomplished using the complementary branch followed by an unconditional jump:

```
BEQ x1, x30, 8  
J    far_addr
```

A misaligned instruction fetch exception will be generated if the address of the branch taken is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

CSRRW

Atomic read/write CSR

Instruction Format: I

Opcode Type: SYSTEM

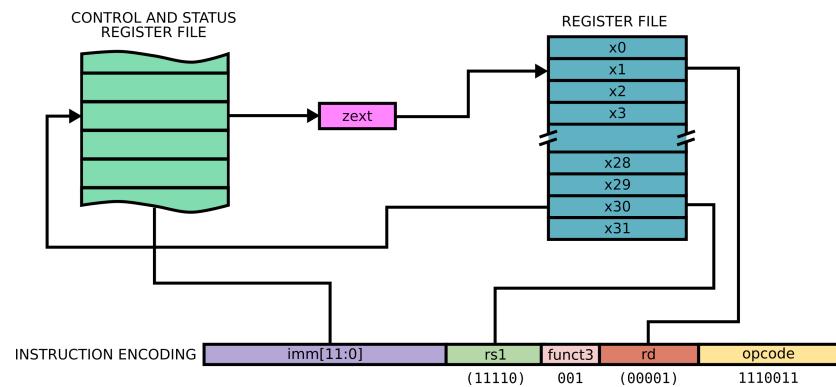
Assembly: CSRRW rd, rs1, csr

Function: $(rd) \leftarrow \text{zext}(\text{csr}[\text{imm}]) \leftarrow (\text{rs1})$

Description:

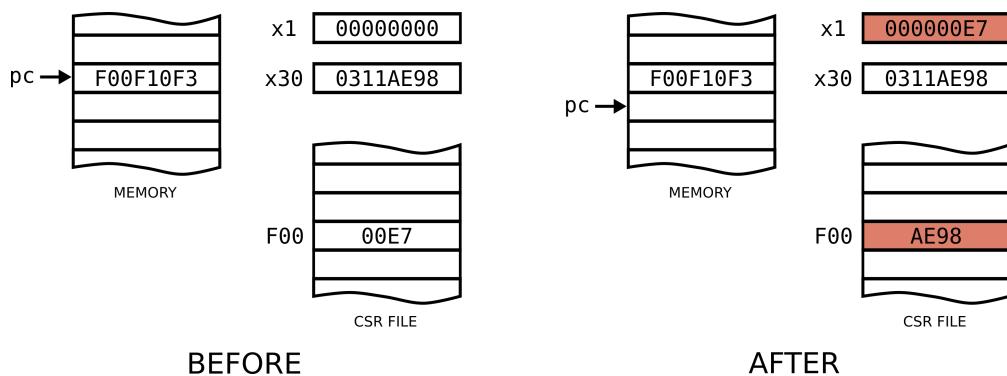
Stores the zero-extended old value of the Control and Status Register in the destination register, and stores the contents of rs1 in the CSR.

Data Path:



Example: CSRRW x1, x30, 0xF00

In this example, CSR 0xF00 is a 16-bit register.



Notes:

If rd and rs1 are the same, this instruction is effectively a swap.

If rd is x0, then the CSR is not read, and so no side-effect of reading the CSR occurs.

Processor implementations may trap SYSTEM opcodes in a software handler, handle them in hardware, or a mix of both.

JAL

Jump and link

Instruction Format: J

Opcode Type: JAL

Assembly: JAL rd, int21e

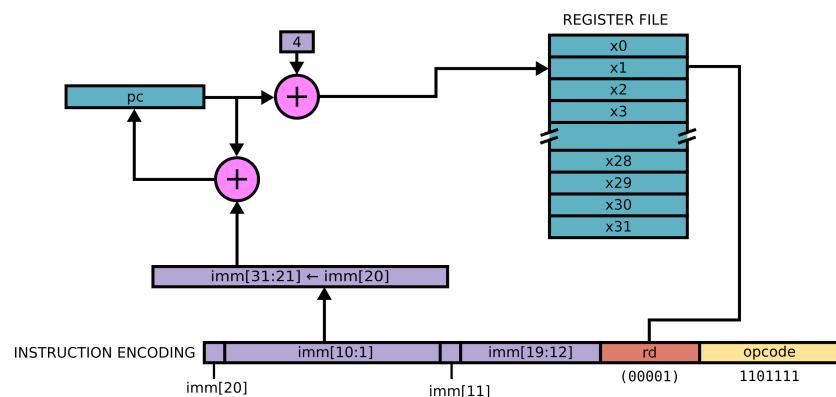
Function:

$$\begin{aligned}(rd) &\leftarrow (pc) + 4 \\ (pc) &\leftarrow (pc) + \text{sext}(imm)\end{aligned}$$

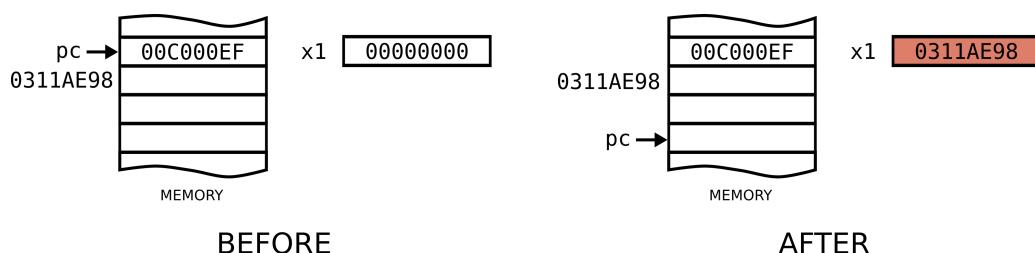
Description:

The address of the next instruction is written to rd, and the signed offset is added to the current program counter.

Data Path:



Example: JAL x1, 12



Notes:

Offsets are limited to even addresses in the range [-1048576, 1048574]. Jumps to relative addresses outside this range can be accomplished using an extra register and the AUIPC and JALR instructions. For example, jumping to pc+0x76543210 can be implemented as follows:

```
AUIPC x30, 0x76543000  
JALR  x1, 0x20C(x30)
```

A misaligned instruction fetch exception will be generated if the target address is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

JALR

Jump and link register

Instruction Format: I

Opcode Type: JALR

Assembly: JALR rd, int12(rs1)

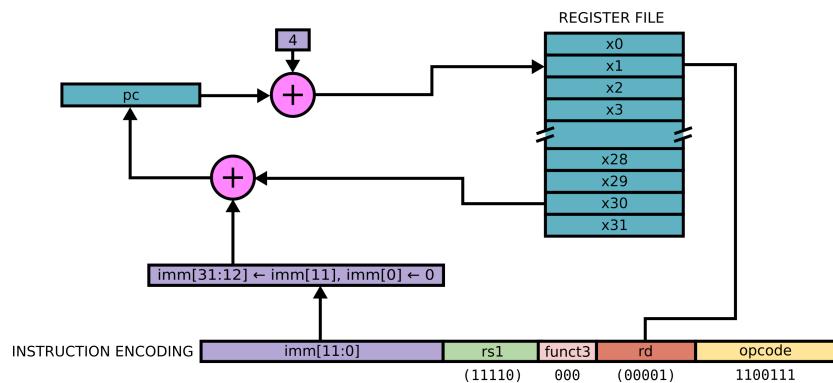
Function:

$$\begin{aligned}(rd) &\leftarrow (pc) + 4 \\ (pc) &\leftarrow (pc) + ((rs1) + \text{sext}(imm)) \wedge \bar{I}\end{aligned}$$

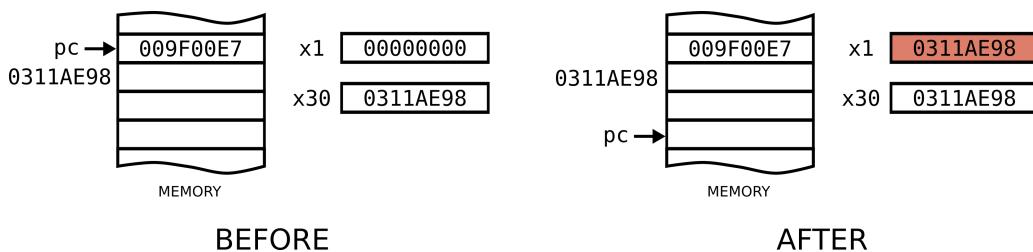
Description:

The address of the next instruction is written to rd. The signed offset is added to the contents of rs1, and the result's least significant bit is set to zero. Then the result is loaded into the program counter.

Data Path:



Example: JALR x1, 9(x30)



Notes:

Offsets are limited to even addresses in the range [-2048, 2046]. Jumps to relative indirect addresses outside this range can be accomplished using an extra register. For example, jumping to 0x76543210(x30) can be implemented as follows:

```
LUI    x29, 0x76543  
ADD    x29, x29, x30  
JALR   x1, 0x210(x29)
```

A misaligned instruction fetch exception will be generated if the target address is not aligned on a 32-bit boundary, except in processors that support extensions with 16-bit aligned instructions, such as the compressed instruction set extension C.

LB

Load byte

Instruction Format: I

Opcode Type: LOAD

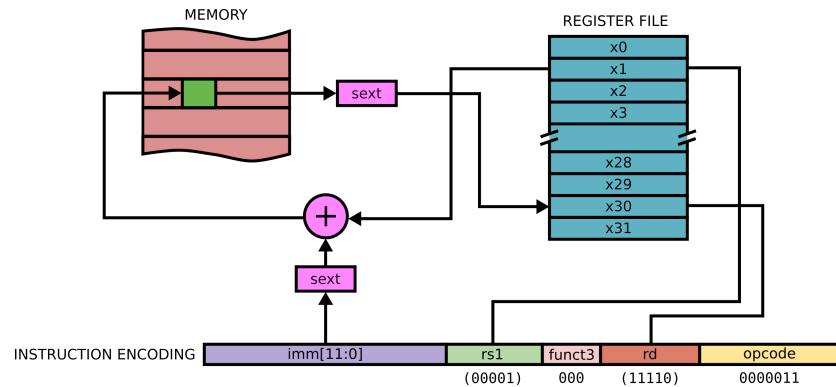
Assembly: LB rd, int12(rs1)

Function: $(rd) \leftarrow \text{sext}([(rs1) + \text{sext}(imm)]_b)$

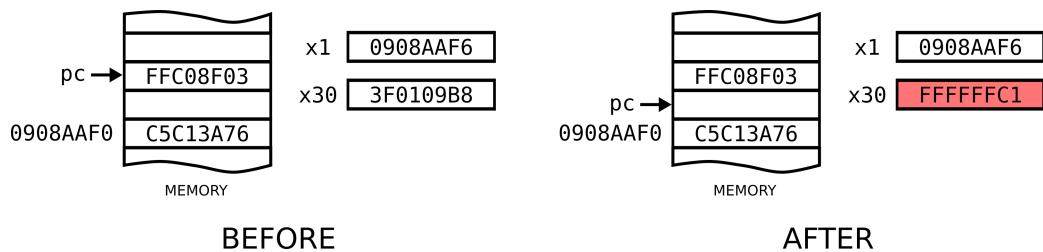
Description:

Adds the signed offset to the contents of register rs1, treating it as a memory address. Retrieves a byte from memory at that address, and sign-extends it. The result is placed in the destination register.

Data Path:



Example: LB x30, -4(x1)



In this example, the little-endian layout in memory at address 0908AAF0 is 76 3A C1 C5. The address to access is 0908AAF2, which contains the byte C1. When sign-extended to 32 bits, this is FFFFFFFC1.

LBU

Load byte unsigned

Instruction Format: I

Opcode Type: LOAD

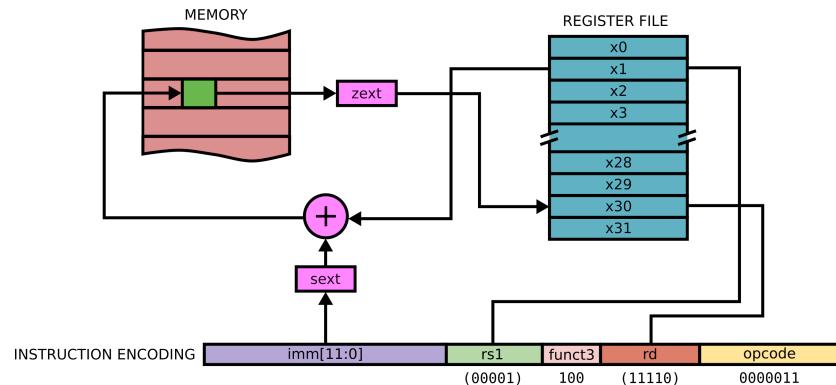
Assembly: LBU rd, int12(rs1)

Function: $(rd) \leftarrow \text{zext}([(rs1) + \text{sext}(imm)]_h)$

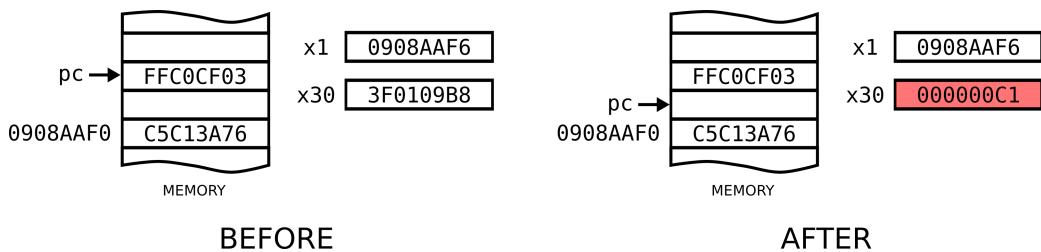
Description:

Adds the signed offset to the contents of register rs1, treating it as a memory address. Retrieves a byte from memory at that address, and zero-extends it. The result is placed in the destination register.

Data Path:



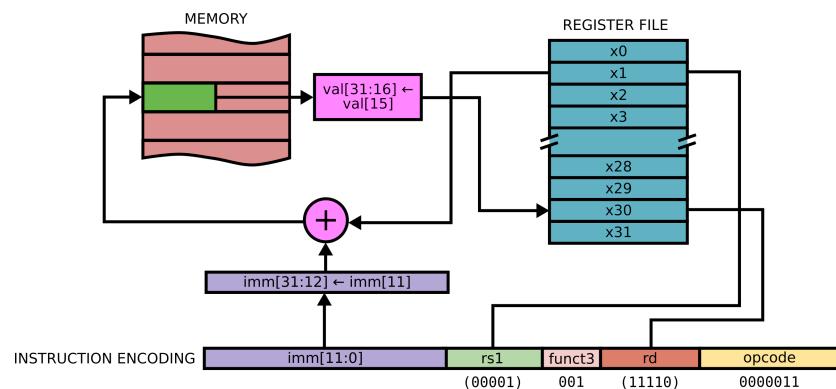
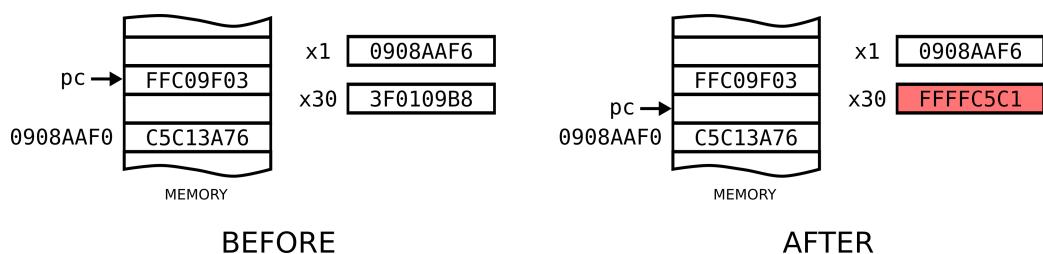
Example: LBU x30, -4(x1)



In this example, the little-endian layout in memory at address 0908AAF0 is 76 3A C1 C5. The address to access is 0908AAF2, which contains the byte C1. When zero-extended to 32 bits, this is 000000C1.

Instruction Format: I**Opcode Type:** LOAD**Assembly:** LH rd, int12(rs1)**Function:** $(rd) \leftarrow \text{sext}([(rs1) + \text{sext}(imm)]_h)$ **Description:**

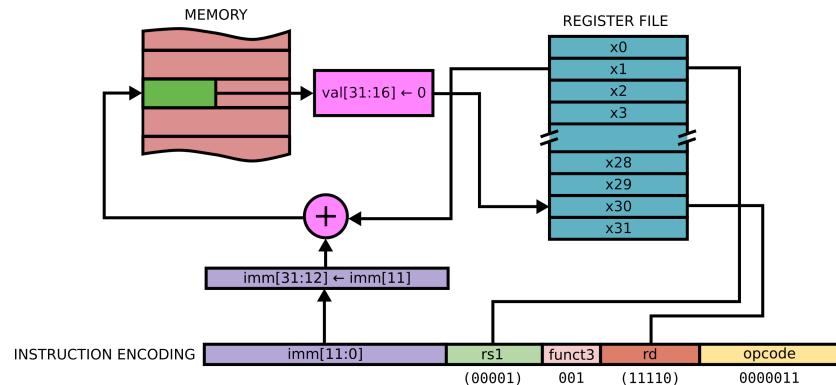
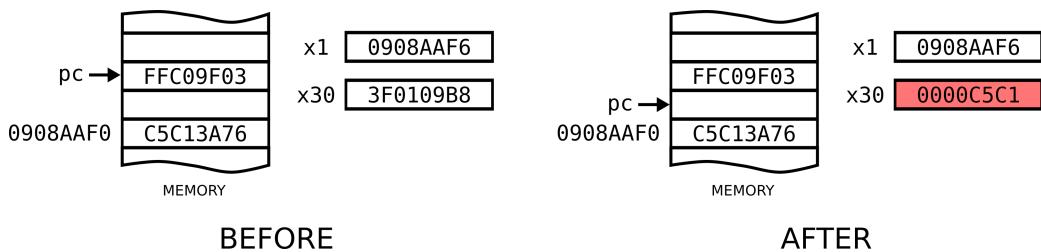
Adds the signed offset to the contents of register rs1, treating it as a memory address. Retrieves a little-endian 16-bit value from memory at that address, and sign-extends it. The result is placed in the destination register.

Data Path:**Example:** LH x30, -4(x1)

In this example, the little-endian layout in memory at address 0908AAF0 is 76 3A C1 C5. The address to access is 0908AAF2, which contains the bytes C1 C5. The corresponding little-endian value is C5C1. When sign-extended to 32 bits, this is FFFFC5C1.

Instruction Format: I**Opcode Type:** LOAD**Assembly:** LHU rd, int12(rs1)**Function:** $(rd) \leftarrow \text{zext}([(rs1) + \text{sext}(imm)]_h)$ **Description:**

Adds the signed offset to the contents of register rs1, treating it as a memory address. Retrieves a little-endian 16-bit value from memory at that address, and zero-extends it. The result is placed in the destination register.

Data Path:**Example:** LHU x30, -4(x1)

In this example, the little-endian layout in memory at address 0908AAF0 is 76 3A C1 C5. The address to access is 0908AAF2, which contains the bytes C1 C5. The corresponding little-endian value is C5C1. When zero-extended to 32 bits, this is 0000C5C1.

LUI

Load upper immediate

Instruction Format: U

Opcode Type: LUI

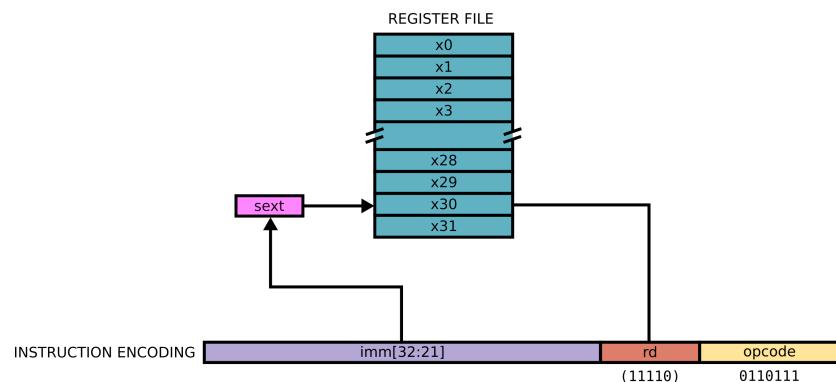
Assembly: LUI rd, uint20

Function: $(rd) \leftarrow \text{sext}(\text{imm} \ll 12)$

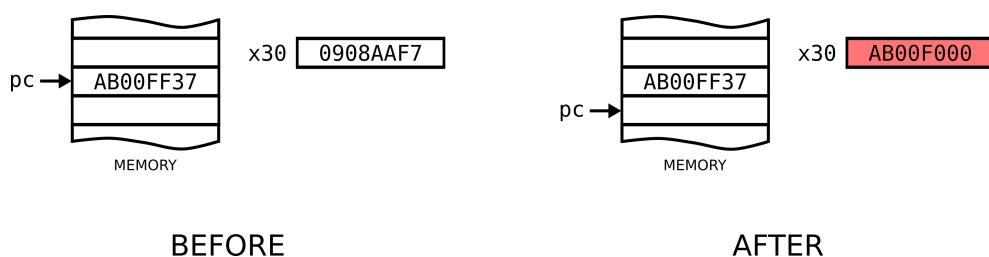
Description:

Constructs a 32-bit integer whose upper 20 bits are the immediate value, filling the lower 12 bits with zeros. The result is then sign-extended to XLEN bits and placed in the destination register.

Data Path:



Example: LUI x30, 0xAB00F



Notes:

The LUI instruction is used to construct constants outside the 12-bit range.

LW

Load word

Instruction Format: I

Opcode Type: LOAD

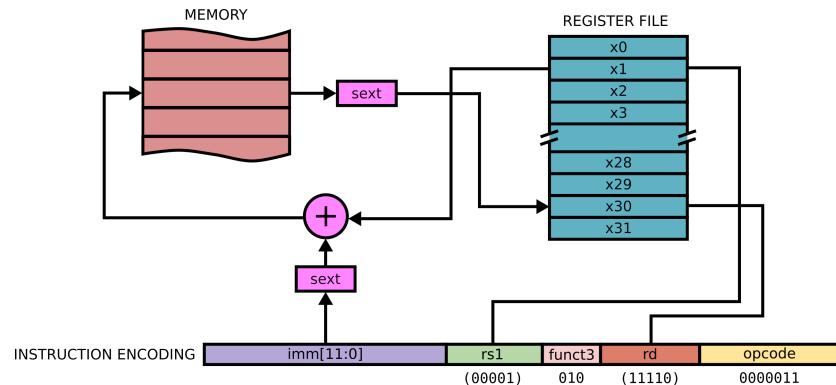
Assembly: LW rd, int12(rs1)

Function: $(rd) \leftarrow \text{sext}([(rs1) + \text{sext}(imm)]_w)$

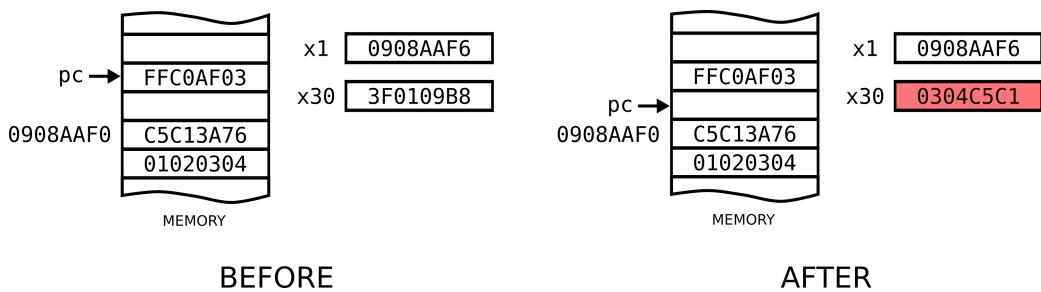
Description:

Adds the signed offset to the contents of register rs1, treating it as a memory address. Retrieves a little-endian 32-bit value from memory at that address, and sign-extends it. The result is placed in the destination register.

Data Path:



Example: LW x30, -4(x1)



In this example, the little-endian layout in memory at address 0908AAF0 is 76 3A C1 C5 04 03 02 01. The address to access is 0908AAF2, which contains the bytes C1 C5 04 03. The corresponding little-endian value is 0304C5C1.

NOP

No operation

Instruction Format: I

Opcode Type: LOAD

Assembly: (*pseudo-instruction*) NOP

Function: $(pc) \leftarrow (pc) + 4$

Description:

Does not change any user-visible state, except for advancing the program counter.

Notes:

While there are many instructions that can do nothing, the standard NOP is implemented as an ADDI instruction:

ADDI x0, x0, 0

OR

Bitwise or (register/register)

Instruction Format: R

Opcode Type: OP

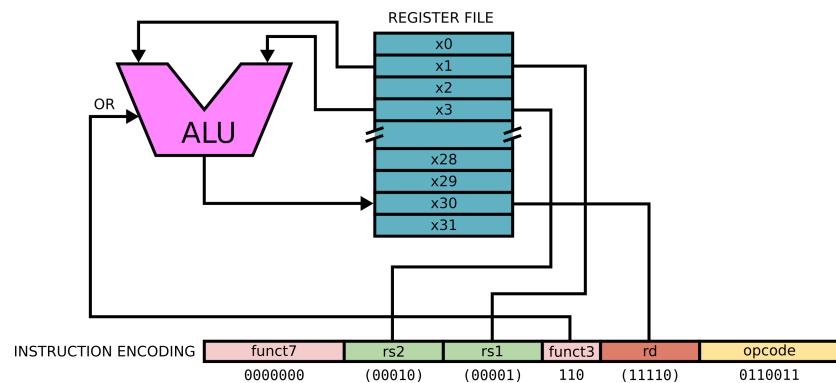
Assembly: OR rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) \vee (rs2)$

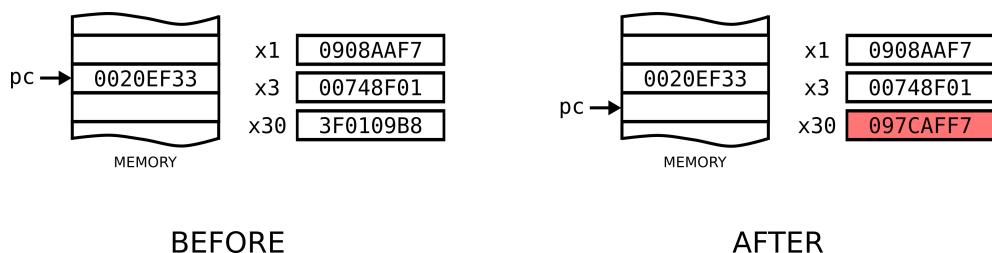
Description:

Performs a bitwise OR between the contents of register rs1 and the contents of register rs2. The result is placed in register rd.

Data Path:



Example: OR x30, x1, x3



ORI

Bitwise or (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

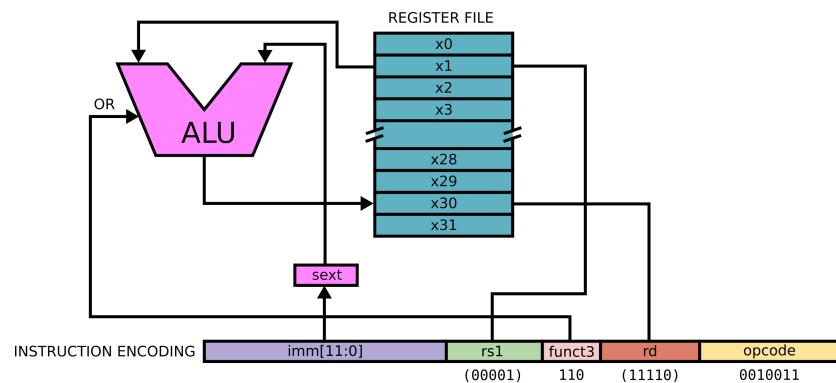
Assembly: OR rd, rs1, int12

Function: $(rd) \leftarrow (rs1) \vee \text{sext(immm)}$

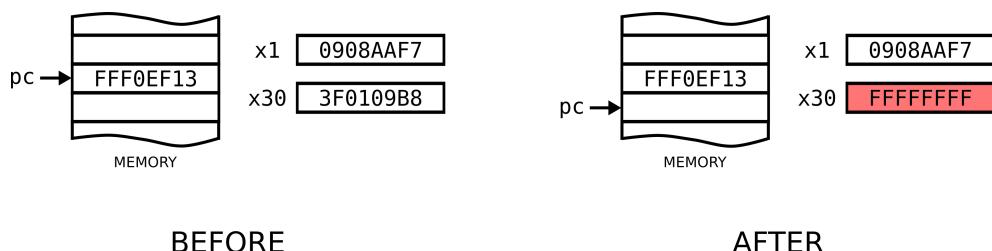
Description:

Performs a bitwise OR between the contents of register rs1 and the sign-extended immediate value. The result is placed in register rd.

Data Path:



Example: ORI x30, x1, -1



SB

Store byte

Instruction Format: S

Opcode Type: STORE

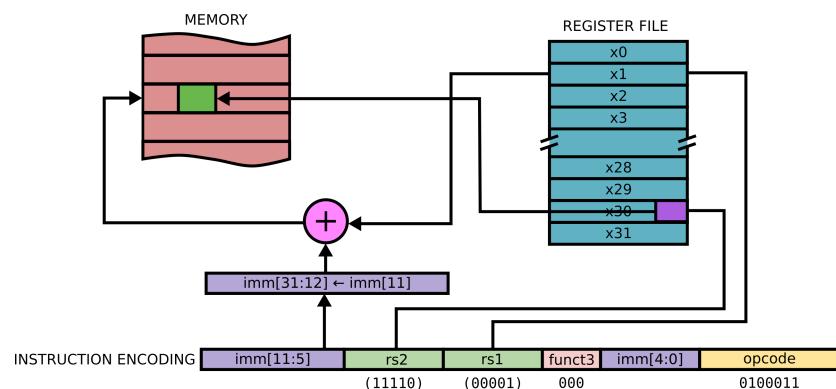
Assembly: SB int12(rs1), rs2

Function: $[(rs1) + \text{sext}(\text{imm})]_b \leftarrow (rs2)_b$

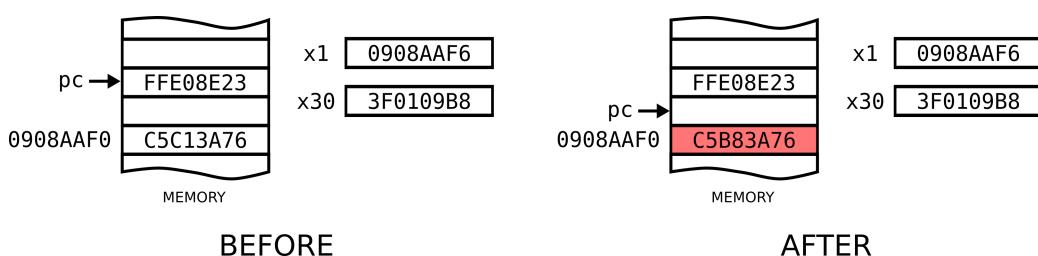
Description:

Adds the signed offset to the contents of register rs1, treating it as a memory address. Stores the bottom 8 bits of rs2 into memory at that address.

Data Path:



Example: SB -4(x1), x30



In this example, the bytes in memory starting at address 0908AAF0 are 76 3A C1 C5. The address to access is 0908AAF2, which contains the byte C1. The value of the lower 8 bits in x30 is B8. This is stored in memory, replacing C1, so that now the bytes in memory starting at address 0908AAF0 are 76 3A B8 C5.

SH

Store half word

Instruction Format: S

Opcode Type: STORE

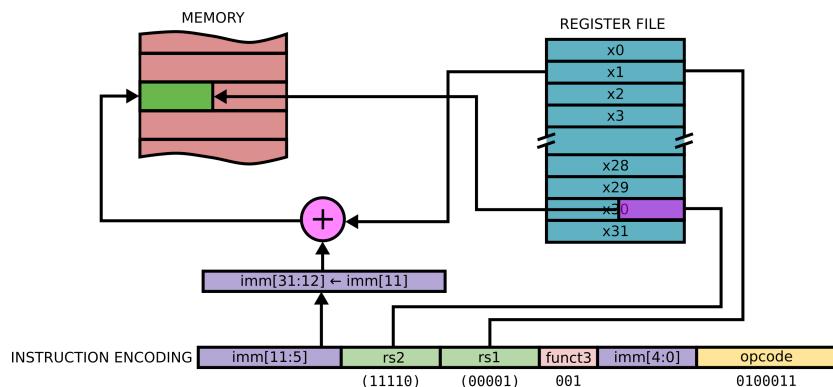
Assembly: SH int12(rs1), rs2

Function: $[(rs1) + \text{sext}(\text{imm})]_h \leftarrow (rs2)_h$

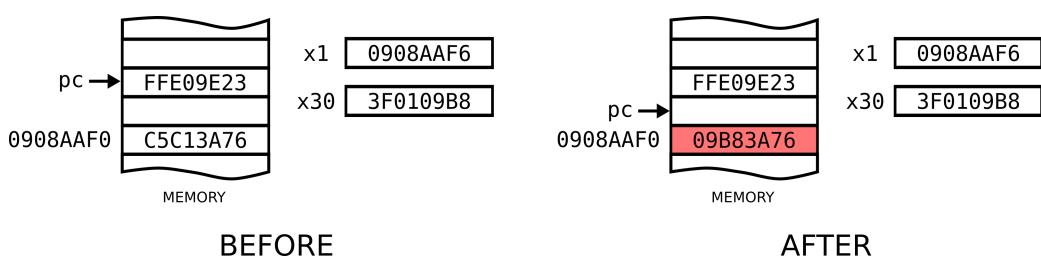
Description:

Adds the signed offset to the contents of register rs1, treating it as a memory address. Stores the bottom 16 bits of rs2 into memory at that address.

Data Path:



Example: SH -4(x1), x30



In this example, the bytes in memory starting at address 0908AAF0 are 76 3A C1 C5. The address to access is 0908AAF2, which contains the bytes C1 C5. The value of the lower 8 bits in x30 is 09B8. This is stored in memory as the little-endian bytes B8 09, replacing C1 C5, so that now the bytes in memory starting at address 0908AAF0 are 76 3A B8 C5.

SLL

Shift left logical (register/register)

Instruction Format: R

Opcode Type: OP

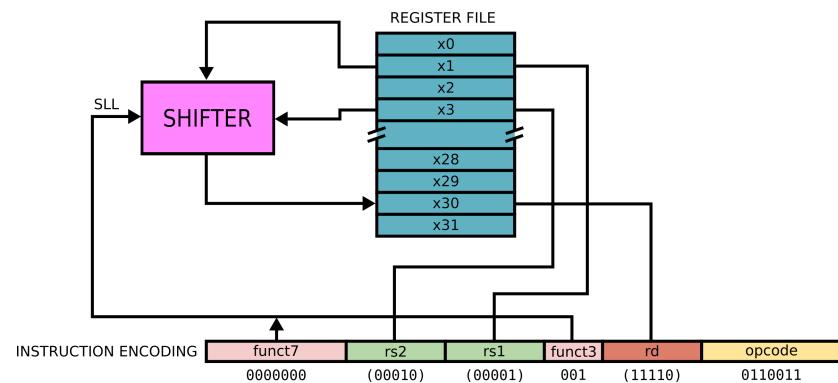
Assembly: SLL rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) \ll (rs2)_n$

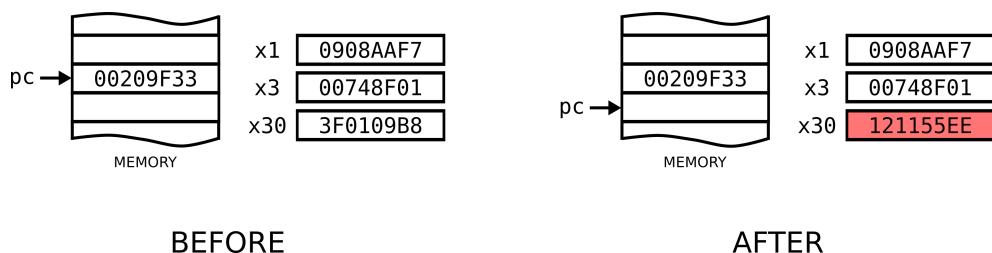
Description:

Shifts the contents of register rs1 left by the lowest $\log_2(XLEN)$ bits of the contents of register rs2, shifting in zeroes. The result is stored in the destination register.

Data Path:



Example: SLL x30, x1, x3



SLLI

Shift left logical (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

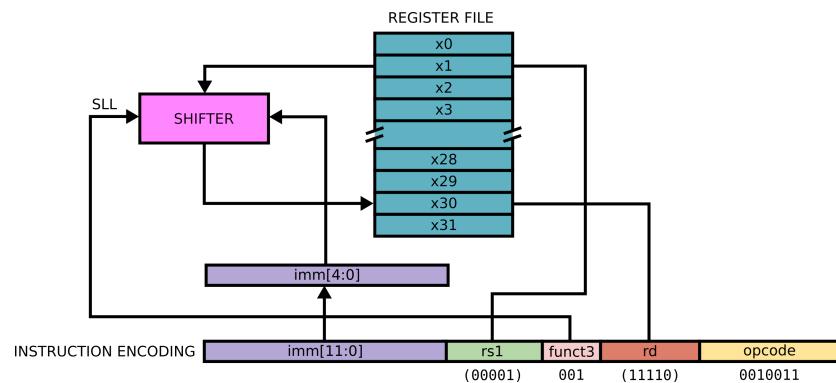
Assembly: SLLI rd, rs1, uint12

Function: $(rd) \leftarrow (rs1) \ll imm_n$

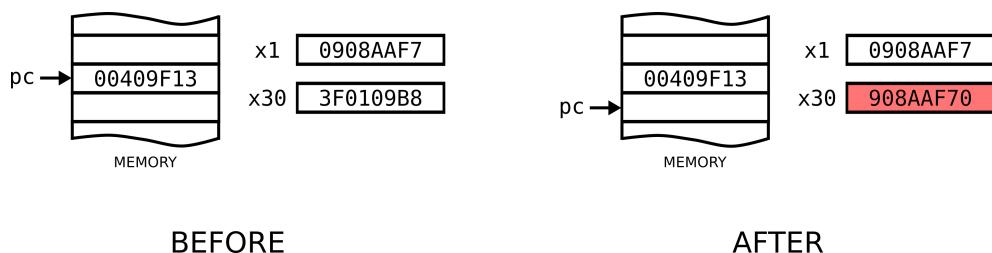
Description:

Shifts the contents of register rs1 left by the lowest $\log_2(XLEN)$ bits of the immediate number, shifting in zeroes. The result is stored in the destination register.

Data Path:



Example: SLLI x30, x1, 4



SLT

Set if less than (register/register)

Instruction Format: R

Opcode Type: OP

Assembly: SLT rd, rs1, rs2

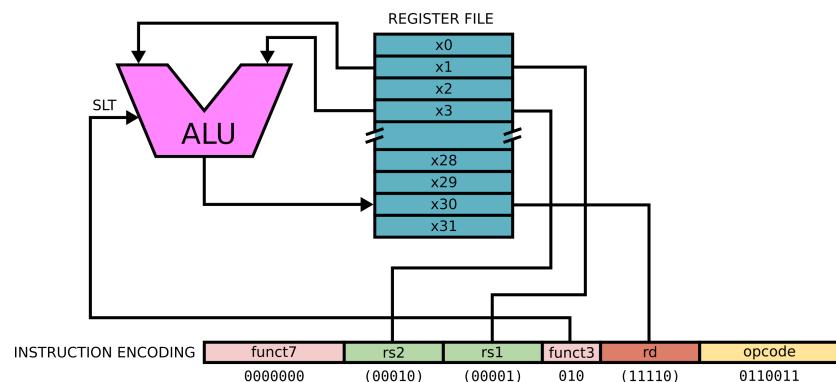
Function:

$$(rd) \leftarrow \begin{cases} 1 & \text{if } (rs1) < (rs2) \\ 0 & \text{otherwise} \end{cases}$$

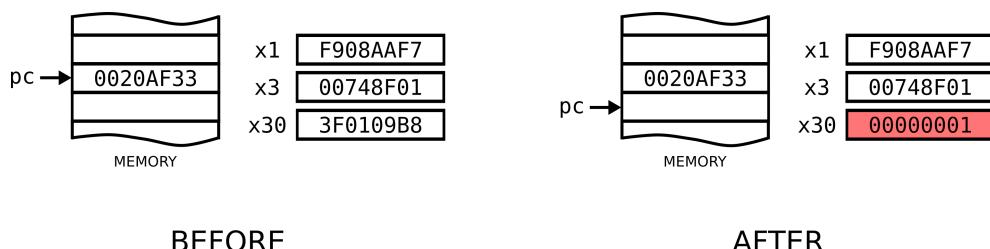
Description:

Sets the destination register to 1 if the contents of register rs1 is less than the contents of register rs2 in a signed comparison, otherwise 0.

Data Path:



Example: SLT x30, x1, x3



SLTI

Set if less than (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

Assembly: SLTI rd, rs1, int12

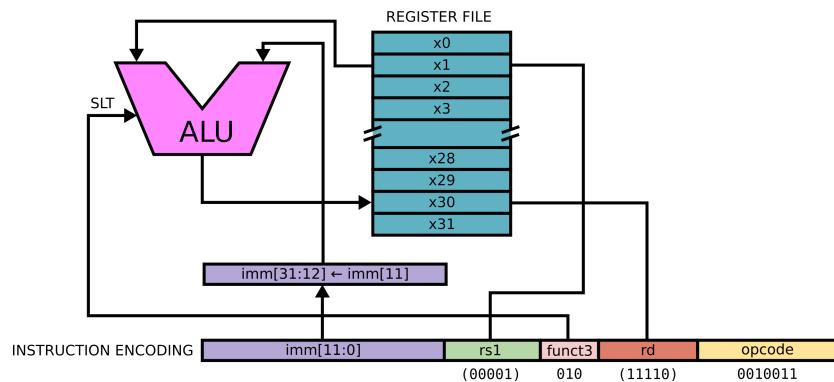
Function:

$$(rd) \leftarrow \begin{cases} 1 & \text{if } (rs1) < \text{sext(imm)} \\ 0 & \text{otherwise} \end{cases}$$

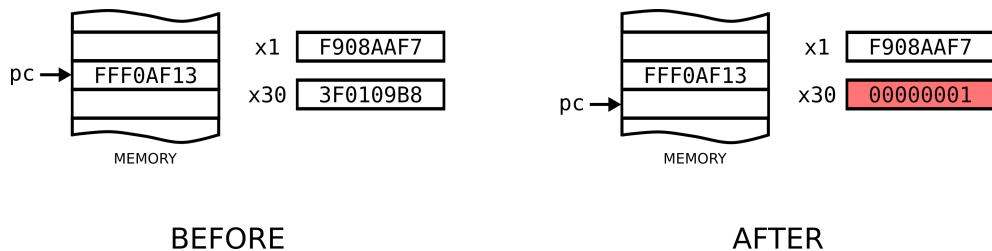
Description:

Sets the destination register to 1 if the contents of register rs1 is less than the sign-extended immediate value in a signed comparison, otherwise 0.

Data Path:



Example: SLTI x30, x1, -1



Notes:

Overflows of the ADD instruction for signed arithmetic can be checked using SLTI and SLT:

ADD x1, x2, x3
SLTI x4, x3, 0
SLT x5, x1, x2
BNE x4, x5, overflow

SLTIU

Set if less than, unsigned (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

Assembly: SLTIU rd, rs1, uint12

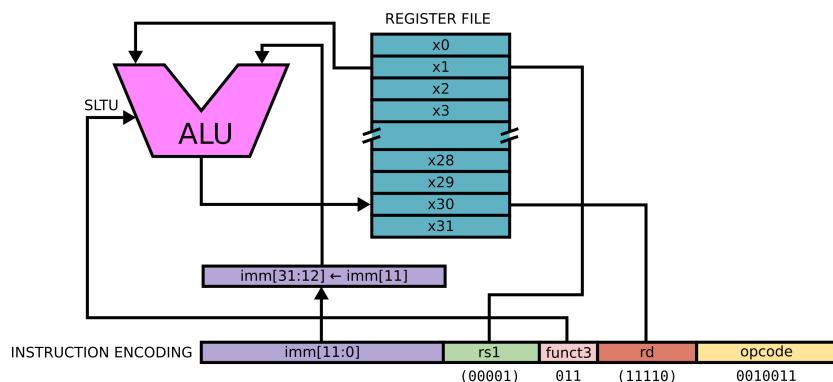
Function:

$$(rd) \leftarrow \begin{cases} 1 & \text{if } (rs1)_u < \text{zext(imm)} \\ 0 & \text{otherwise} \end{cases}$$

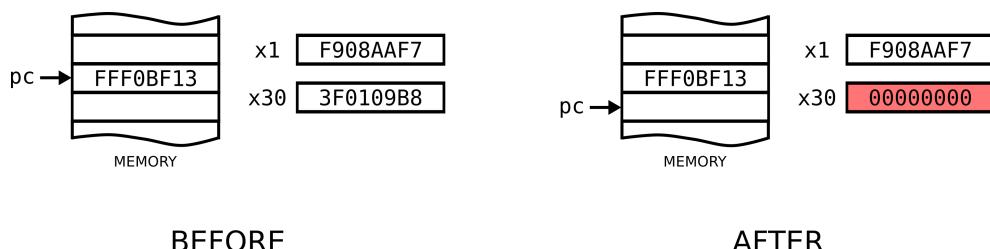
Description:

Sets the destination register to 1 if the contents of register rs1 is less than the zero-extended immediate value in an unsigned comparison, otherwise 0.

Data Path:



Example: SLTIU x30, x1, 0xFFFF



SLTU

Set if less than, unsigned (register/register)

Instruction Format: R

Opcode Type: OP

Assembly: SLTU rd, rs1, rs2

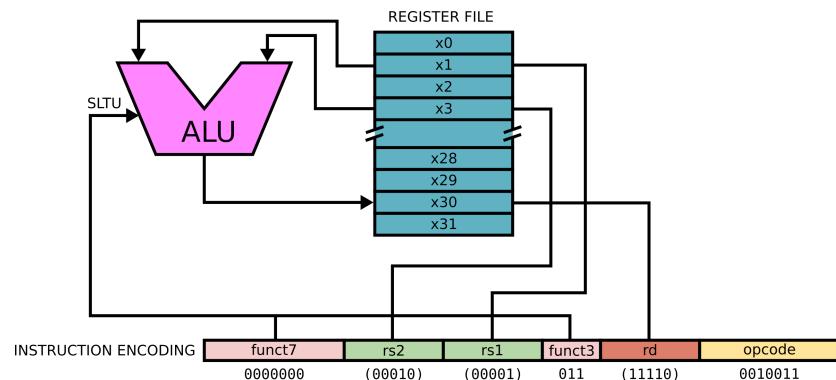
Function:

$$(rd) \leftarrow \begin{cases} 1 & \text{if } (rs1)_u < (rs2)_u \\ 0 & \text{otherwise} \end{cases}$$

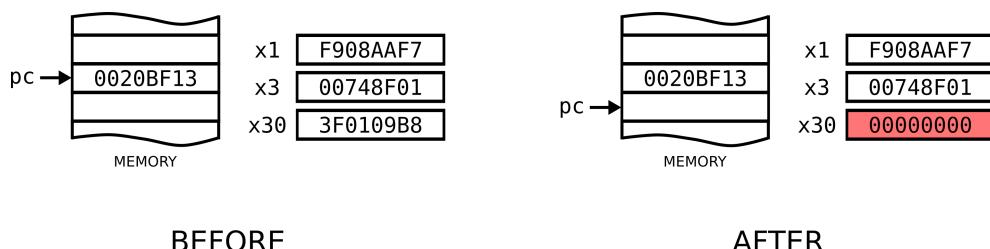
Description:

Sets the destination register to 1 if the contents of register rs1 is less than the contents of register rs2 in an unsigned comparison, otherwise 0.

Data Path:



Example: SLTU x30, x1, x3



SRA

Shift right arithmetic (register/register)

Instruction Format: R

Opcode Type: OP

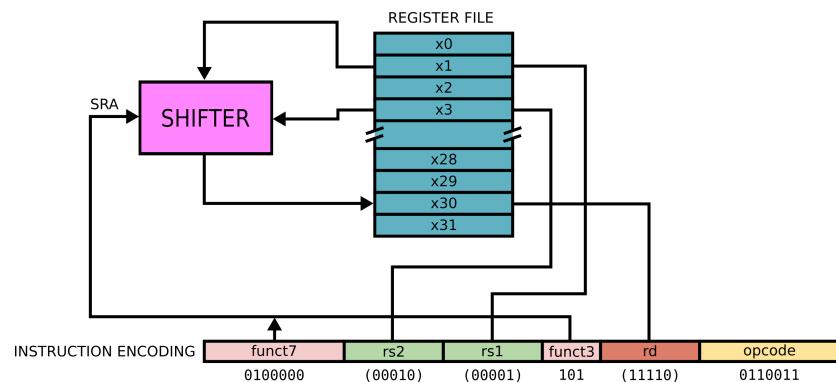
Assembly: SRA rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) \gg (rs2)_n$

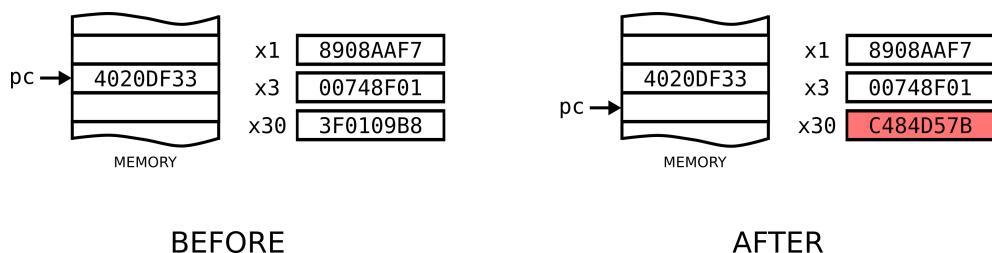
Description:

Shifts the contents of register rs1 right by the lowest $\log_2(XLEN)$ bits of the contents of register rs2, shifting in the most significant bit of register rs1. The result is stored in the destination register.

Data Path:



Example: SRA x30, x1, x3



SRAI

Shift right arithmetic (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

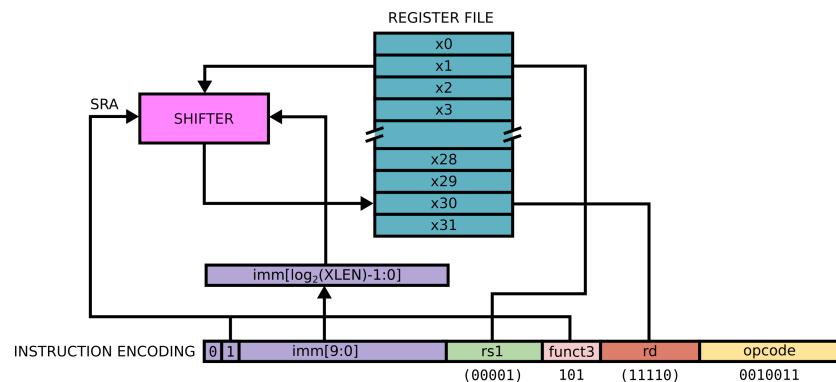
Assembly: SRAI rd, rs1, uintN

Function: $(rd) \leftarrow (rs1) \gg (imm)_n$

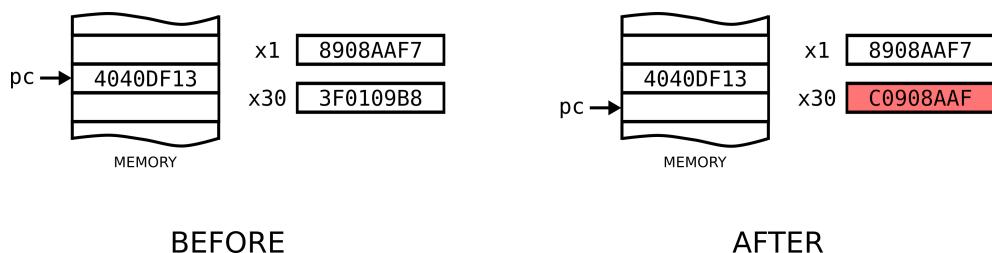
Description:

Shifts the contents of register rs1 right by the lowest $\log_2(XLEN)$ bits of the immediate value, shifting in the most significant bit of register rs1. The result is stored in the destination register.

Data Path:



Example: SRAI x30, x1, 4



SRL

Shift right logical (register/register)

Instruction Format: R

Opcode Type: OP

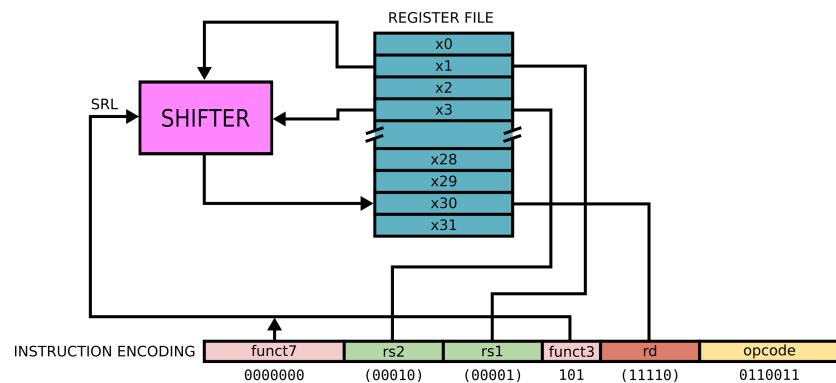
Assembly: SRL rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) \ggg (rs2)_n$

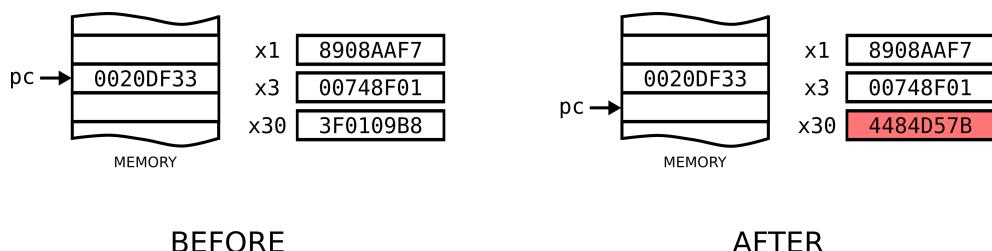
Description:

Shifts the contents of register rs1 right by the lowest $\log_2(XLEN)$ bits of the contents of register rs2, shifting in zeroes. The result is stored in the destination register.

Data Path:



Example: SRL x30, x1, x3



SRLI

Shift right logical (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

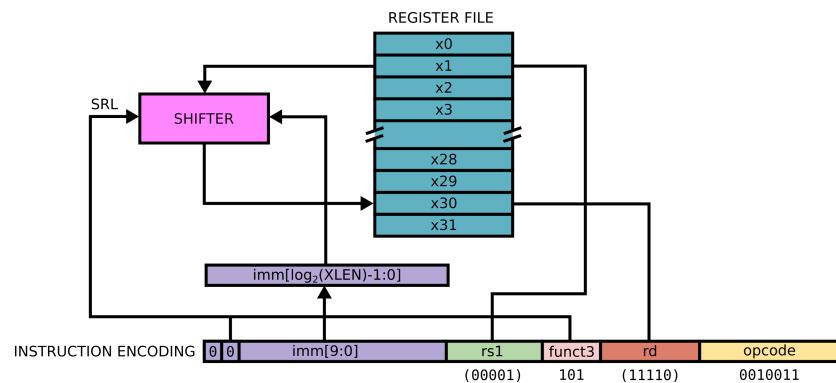
Assembly: SRLI rd, rs1, uintN

Function: $(rd) \leftarrow (rs1) \ggg (imm)_n$

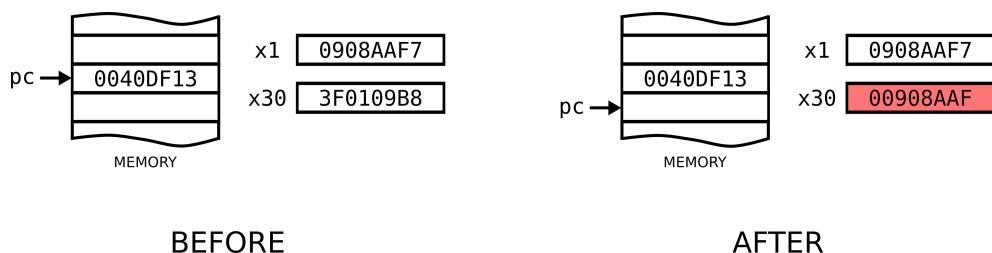
Description:

Shifts the contents of register rs1 right by the lowest $\log_2(XLEN)$ bits of the immediate value, shifting in zeroes. The result is stored in the destination register.

Data Path:



Example: SRLI x30, x1, 4



SUB

Subtract (register/register)

Instruction Format: R

Opcode Type: OP

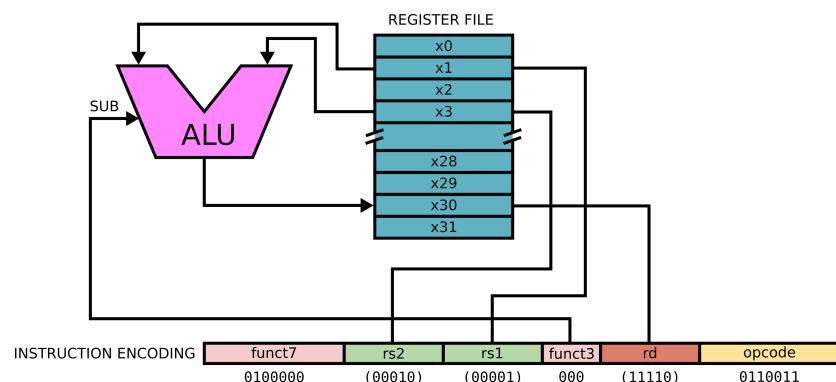
Assembly: SUB rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) - (rs2)$

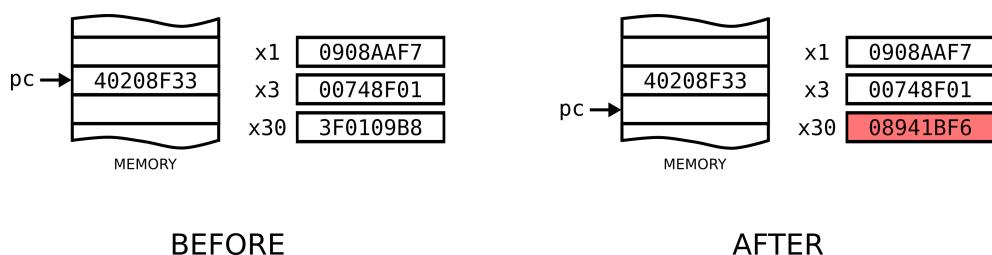
Description:

Subtracts contents of register rs2 from the contents of register rs1. The result is stored in the destination register.

Data Path:



Example: SUB x30, x1, x3



Notes:

There is no separate SUBI instruction. However, SUBI is a pseudo-instruction encoded as ADDI where the immediate value is negated. Thus, the following two instructions are encoded the same way:

SUBI x30, x1, 1
ADDI x30, x1, -1

SW

Store word

Instruction Format: S

Opcode Type: STORE

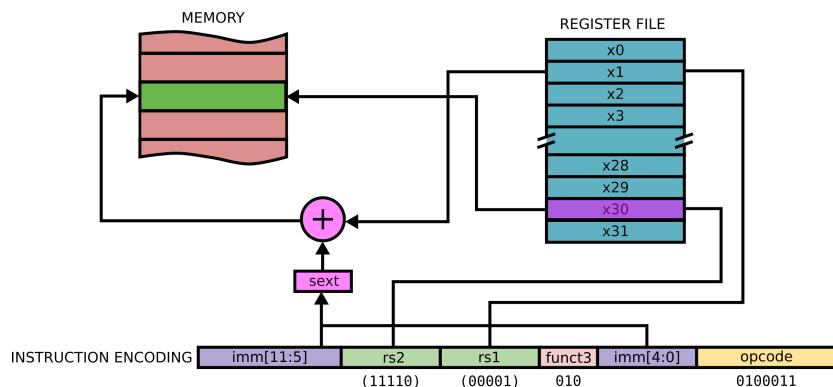
Assembly: SW int12(rs1), rs2

Function: $[(rs1) + \text{sext}(\text{imm})]_w \leftarrow (rs2)_w$

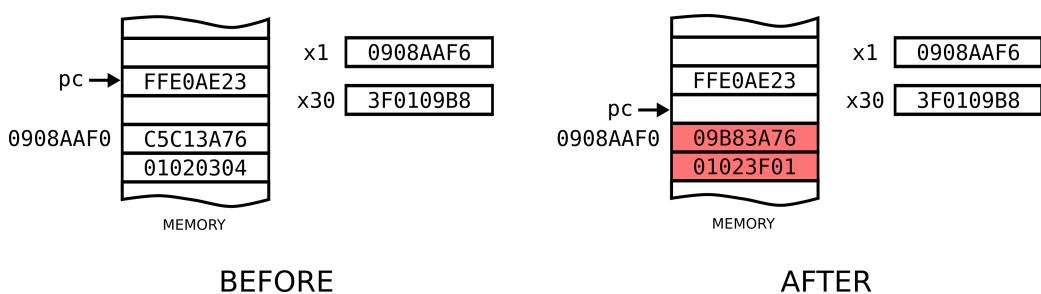
Description:

Adds the signed offset to the contents of register rs1, treating it as a memory address. Stores the bottom 32 bits of rs2 into memory at that address.

Data Path:



Example: SW -4(x1), x30



In this example, the bytes in memory starting at address 0908AAF0 are 76 3A C1 C5 04 03 02 01. The address to access is 0908AAF2, which contains the bytes C1 C5 04 03. The contents of x30 is 3F0109B8. This is stored in memory as the little-endian bytes B8 09 01 3F, replacing C1 C5 04 03, so that now the bytes in memory starting at address 0908AAF0 are 76 3A B8 09 01 3F 02 01.

XOR

Bitwise exclusive or (register/register)

Instruction Format: R

Opcode Type: OP

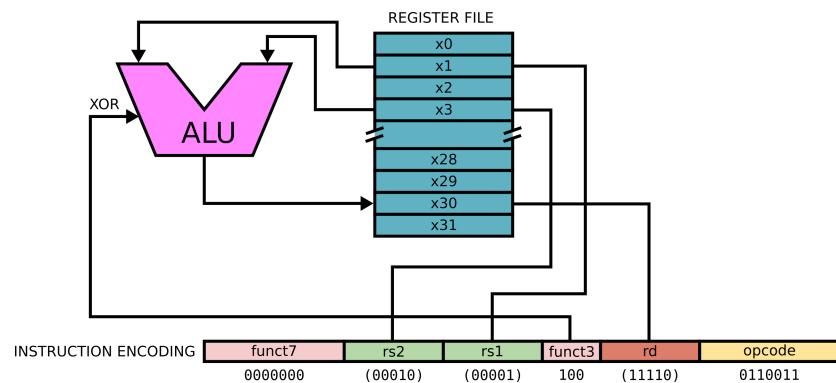
Assembly: XOR rd, rs1, rs2

Function: $(rd) \leftarrow (rs1) \oplus (rs2)$

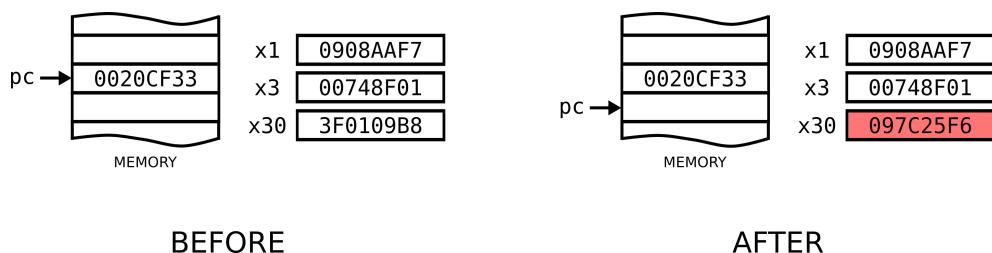
Description:

Performs a bitwise XOR between the contents of register rs1 and the contents of register rs2. The result is placed in register rd.

Data Path:



Example: XOR x30, x1, x3



XORI

Bitwise exclusive or (register/immediate)

Instruction Format: I

Opcode Type: OP-IMM

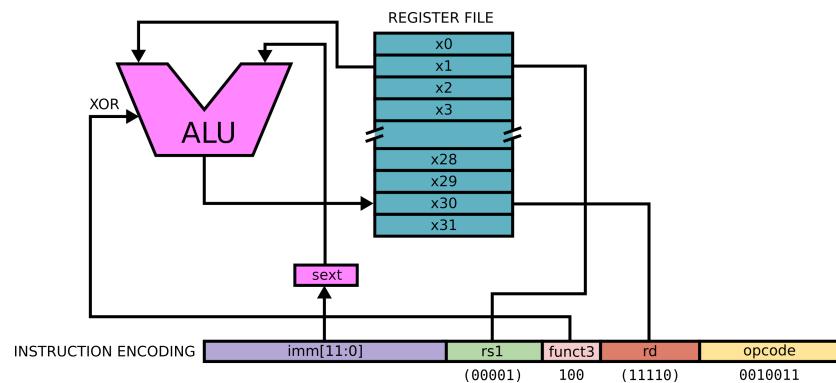
Assembly: XOR rd, rs1, int12

Function: $(rd) \leftarrow (rs1) \oplus \text{sext}(\text{imm})$

Description:

Performs a bitwise XOR between the contents of register rs1 and the sign-extended immediate value. The result is placed in register rd.

Data Path:



Example: X0RI x30, x1, -1

